

# **MPI Collectives**

Computational Science II (CAAM 520)

---

Christopher Thiele

Rice University, Spring 2021

## What are collectives?

So far we have introduced MPI communicators as well as `MPI_Send()`, `MPI_Recv()`, and `MPI_Sendrecv()` to pass messages between **individual** ranks within a communicator.

**Collectives** are operations that are performed by **all** ranks in a communicator.

Collectives provide convenient and efficient implementations of common communication patterns.

## Synchronizing ranks with MPI\_Barrier

We are already familiar with one example of a collective operation: MPI\_Barrier.

```
int MPI_Barrier(MPI_Comm comm)
```

→ Each rank waits for **all** other ranks to reach the barrier.


## Broadcasting data with `MPI_Bcast`

To send data from one rank to all other ranks in the communicator, use `MPI_Bcast`:

```
int MPI_Bcast(void *buf,  
              int count,  
              MPI_Datatype datatype,  
              int root,  
              MPI_Comm comm)
```

→ The rank specified by the `root` argument broadcasts to all other ranks.

# Broadcasting data with `MPI_Bcast`



---

Image source: [https://www.wikiwand.com/de/Message\\_Passing\\_Interface](https://www.wikiwand.com/de/Message_Passing_Interface)


## Scattering data with MPI\_Scatter

MPI\_Scatter works much like MPI\_Bcast, but it sends different data to each rank.

```
int MPI_Scatter(const void *sendbuf,  
               int sendcount,  
               MPI_Datatype sendtype,  
               void *recvbuf,  
               int recvcount,  
               MPI_Datatype recvtype,  
               int root,  
               MPI_Comm comm)
```

→ MPI\_Scatter sends sendcount items to **each** rank.

# Scattering data with `MPI_Scatter`



---

Image source: [https://www.wikiwand.com/de/Message\\_Passing\\_Interface](https://www.wikiwand.com/de/Message_Passing_Interface)

## Scattering data with MPI\_Scatter

If we need to send a different amount of data to each rank, we can use MPI\_Scatterv.

```
int MPI_Scatterv(const void *sendbuf,  
                const int *sendcounts,  
                const int *displs,  
                MPI_Datatype sendtype,  
                void *recvbuf,  
                int recvcnt,  
                MPI_Datatype recvtype,  
                int root,  
                MPI_Comm comm)
```

→ Sends `sendcounts[r]` items to rank `r` starting with `sendbuf[displs[r]]`.



## Gathering data with MPI\_Gather

The inverse operation to MPI\_Scatter is MPI\_Gather. It collects data from all ranks on the root rank.

```
int MPI_Gather(const void *sendbuf,  
              int sendcount,  
              MPI_Datatype sendtype,  
              void *recvbuf,  
              int recvcount,  
              MPI_Datatype recvtype,  
              int root,  
              MPI_Comm comm)
```

→ MPI\_Gather receives `recvcount` items from *each* rank.

# Gathering data with MPI\_Gather




Image source: [https://www.wikiwand.com/de/Message\\_Passing\\_Interface](https://www.wikiwand.com/de/Message_Passing_Interface)

## Gathering data with `MPI_Gather`

To collect a different amount of data from each rank, use `MPI_Gatherv`.

```
int MPI_Gatherv(const void *sendbuf,  
                int sendcount,  
                MPI_Datatype sendtype,  
                void *recvbuf,  
                const int *recvcounts,  
                const int *displs,  
                MPI_Datatype recvtype,  
                int root,  
                MPI_Comm comm)
```

→ The root rank receives `recvcounts[r]` from rank `r` and stores them at `sendbuf + displs[r]`.

## Gathering data with MPI\_Allgather

MPI\_Allgather works just like MPI\_Gather, but **all** ranks gather all data.

```
int MPI_Allgather(const void *sendbuf,
                  int sendcount,
                  MPI_Datatype sendtype,
                  void *recvbuf,
                  int recvcount,
                  MPI_Datatype recvtype,
                  MPI_Comm comm)
```

→ Again, there is a more general MPI\_Allgatherv function, too.

# Gathering data with MPI\_Allgather



Image source: [https://www.wikiwand.com/de/Message\\_Passing\\_Interface](https://www.wikiwand.com/de/Message_Passing_Interface)

## Redistributing data with MPI\_Alltoall

MPI\_Alltoall resembles MPI\_Allgather, but each rank now gathers different data. The action of MPI\_Alltoall is best explained in a picture (next slide).

```
int MPI_Alltoall(const void *sendbuf,  
                int sendcount,  
                MPI_Datatype sendtype,  
                void *recvbuf,  
                int recvcount,  
                MPI_Datatype recvtype,  
                MPI_Comm comm)
```

# Redistributing data with MPI\_Alltoall




Image source: [https://www.wikiwand.com/de/Message\\_Passing\\_Interface](https://www.wikiwand.com/de/Message_Passing_Interface)

## Redistributing data with MPI\_Alltoall

Again, there is a function MPI\_Alltoallv that can be used if each rank receives a different amounts of data.

```
int MPI_Alltoallv(const void *sendbuf,
                  const int *sendcounts,
                  const int *sdispls,
                  MPI_Datatype sendtype,
                  void *recvbuf,
                  const int *recvcounts,
                  const int *rdispls,
                  MPI_Datatype recvtype,
                  MPI_Comm comm)
```



# Reductions

Some collective operations perform computations in addition to message passing.

`MPI_Reduce` works much like `MPI_Gather`, but the gathered data is combined using an operator.

```
int MPI_Reduce(const void *sendbuf,  
               void *recvbuf,  
               int count,  
               MPI_Datatype datatype,  
               MPI_Op op,  
               int root,  
               MPI_Comm comm)
```

→ Compare this to OpenMP's `reduction` clause.

# Reductions

Possible operations include

```
typedef enum {  
    MPI_MAX,  
    MPI_MIN,  
    MPI_SUM,  
    MPI_PROD,  
    MPI_REPLACE,  
    // etc.  
} MPI_Op;
```

→ Users can define their own operations, too.

# Reductions

MPI\_Allreduce works like MPI\_Reduce, but now *every* rank performs a reduction.

```
int MPI_Allreduce(const void *sendbuf,
                  void *recvbuf,
                  int count,
                  MPI_Datatype datatype,
                  MPI_Op op,
                  MPI_Comm comm)
```

# Reductions

**Example:** Compute the Euclidean norm of a *distributed* vector.

```
double norm2(const double *x, int n_local)
{
    double sum, sum_local = 0.0;

    for (int i = 0; i < n_local; i++) {
        sum_local += x[i]*x[i];
    }
    MPI_Allreduce(&sum_local, &sum, 1, MPI_DOUBLE,
                  MPI_SUM, MPI_COMM_WORLD);

    return sqrt(sum);
}
```

## Collectives and deadlocks

Collective operations **must** be called by all ranks in the communicator. Otherwise, the collective operation results in a deadlock!

This can be trickier than it seems:

```
compute_vector(x, n_local);  
  
if (rank == 0) {  
    // Deadlock!  
    printf("||x|| = %e\n", norm2(x, n_local));  
}
```


# Collectives and performance

**Question:** Is the code below a good implementation of MPI\_Bcast?

```
if (rank == root) {
    for (int r = 0; r < size; r++) {
        if (r == rank) continue;
        MPI_Send(buffer, count, datatype,
                 r, 999, comm);
    }
}
else {
    MPI_Recv(buffer, count, datatype,
             root, 999, comm, MPI_STATUS_IGNORE);
}
// ...
```

# Collectives and performance

**Answer:** No! We can accelerate the broadcast using a tree structure.



---

Image source: <https://mpitutorial.com/tutorials/mpi-broadcast-and-collective-communication/>